



Experiment – 3

Aim: Implement the following file management tasks in Hadoop:

- Adding files and directories
- Retrieving files
- Deleting files

Hint: A typical Hadoop workflow creates data files (such as log files) elsewhere and copies them into HDFS using one of the above command line utilities.

File Management tasks in Hadoop

1. Create a directory in HDFS at given path(s).

Usage:

```
hadoop fs -mkdir <paths>
```

Example:

```
hadoop fs -mkdir /user/saurzcode/dir1 /user/saurzcode/dir2
```

2. List the contents of a directory.

Usage :

```
hadoop fs -ls <args>
```

Example:

```
hadoop fs -ls /user/saurzcode
```

3. Upload and download a file in HDFS.

Upload:

hadoop fs -put:

Copy single src file, or multiple src files from local file system to the Hadoop data file system

Usage:

```
hadoop fs -put <localsrc> ... <HDFS_dest_Path>
```

Example:

```
hadoop fs -put /home/saurzcode/Samplefile.txt /user/ saurzcode/dir3/
```

Download:

hadoop fs -get:

Copies/Downloads files to the local file system

Usage:

```
hadoop fs -get <hdfs_src> <localdst>
```

Example:

```
hadoop fs -get /user/saurzcode/dir3/Samplefile.txt /home/
```



4. See contents of a file Same as unix cat command:

Usage:

```
hadoop fs -cat <path[filename]>
```

Example:

```
hadoop fs -cat /user/saurzcode/dir1/abc.txt
```

5. Copy a file from source to destination

This command allows multiple sources as well in which case the destination must be a directory.

Usage:

```
hadoop fs -cp <source> <dest>
```

Example:

```
hadoop fs -cp /user/saurzcode/dir1/abc.txt /user/saurzcode/ dir2
```

6. Copy a file from/To Local file system to HDFS

copyFromLocal

Usage:

```
hadoop fs -copyFromLocal <localsrc> URI
```

Example:

```
hadoop fs -copyFromLocal /home/saurzcode/abc.txt /user/ saurzcode/abc.txt
```

Similar to put command, except that the source is restricted to a local file reference.

copyToLocal

Usage:

```
hadoop fs -copyToLocal [-ignorecrc] [-crc] URI <localdst>
```

Similar to get command, except that the destination is restricted to a local file reference.

7. Move file from source to destination.

Note:- Moving files across filesystem is not permitted.

Usage :

```
hadoop fs -mv <src> <dest>
```

Example:

```
hadoop fs -mv /user/saurzcode/dir1/abc.txt /user/saurzcode/ dir2
```

8. Remove a file or directory in HDFS.

Remove files specified as argument. Deletes directory only when it is empty

Usage :

```
hadoop fs -rm <arg>
```

Example:

```
hadoop fs -rm /user/saurzcode/dir1/abc.txt
```

Recursive version of delete.

Usage :

```
hadoop fs -rmr <arg>
```

Example:

```
hadoop fs -rmr /user/saurzcode/
```

9. Display last few lines of a file. Similar to tail command in Unix.

Usage :

```
hadoop fs -tail <path[filename]>
```

Example:

```
hadoop fs -tail /user/saurzcode/dir1/abc.txt
```

10. Display the aggregate length of a file.

Usage :

```
hadoop fs -du <path>
```

Example:

```
hadoop fs -du /user/saurzcode/dir1/abc.tx
```



Experiment – 4

Aim: Run a basic Word Count Map Reduce program to understand Map Reduce Paradigm.

Word Count Map Reduce program to understand Map Reduce Paradigm

Source code:

```
import java.io.IOException;
import java.util.StringTokenizer;
import
org.apache.hadoop.io.IntWritable;
import
org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import
org.apache.hadoop.mapreduce.Mapper; import
org.apache.hadoop.mapreduce.Reducer;
import
org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.mapreduce.Job;
import
org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import
org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
import
org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import
org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.fs.Path;
public class WordCount
{
public static class Map extends
Mapper<LongWritable,Text,Text,IntWritable> { public void
```



```
map(LongWritable key, Text value, Context context) throws
IOException, InterruptedException {
    String line = value.toString();
    StringTokenizer tokenizer = new
    StringTokenizer(line); while
    (tokenizer.hasMoreTokens()) {
        value.set(tokenizer.nextToken());
        context.write(value, new
        IntWritable(1)); }
    }
}

public static class Reduce extends
Reducer<Text, IntWritable, Text, IntWritable> { public void reduce(Text key,
Iterable<IntWritable> values, Context context) throws
IOException, InterruptedException {
    int sum=0;
    (IntWritable x: values)
    {
        sum+=x.get()
    ; }
    context.write(key, new
    IntWritable(sum)); }
}

public static void main(String[] args) throws Exception {
    Configuration conf= new Configuration();
    Job job = new Job(conf, "My Word Count Program");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(Map.class);
    job.setReducerClass(Reduce.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    job.setInputFormatClass(TextInputFormat.class)
    ;
}
```



```

job.setOutputFormatClass(TextOutputFormat.class);
Path outputPath = new Path(args[1]);
//Configuring the input/output path from the filesystem into
the job FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
//deleting the output path automatically from hdfs so that we don't
have to delete it explicitly
outputPath.getFileSystem(conf).delete(outputPath);
//exiting the job only if the flag value becomes
false System.exit(job.waitForCompletion(true) ? 0 :
1);
}
}

```

The entire MapReduce program can be fundamentally divided into three parts:

- Mapper Phase Code
- Reducer Phase Code
- Driver Code

We will understand the code for each of these three parts sequentially.

Mapper code:

```

public static class Map extends
Mapper<LongWritable,Text,Text,IntWritable>
{
public void map(LongWritable key, Text value, Context context)
throws IOException,InterruptedException {
String line = value.toString();

StringTokenizer tokenizer = new StringTokenizer(line);
while (tokenizer.hasMoreTokens()) {
value.set(tokenizer.nextToken());
context.write(value, new
IntWritable(1)); }
}
}

```



We have created a class Map that extends the class Mapper which is already defined in the MapReduce Framework. We define the data types of input and output key/value pair after the class declaration using angle brackets. Both the input and output of the Mapper is a key/value pair.

Input: The key is nothing but the offset of each line in the text file: LongWritable
The value is each individual line (as shown in the figure at the right): Text

Output: The key is the tokenized words: Text. We have the hardcoded value in our case which is 1: IntWritable. Example – Dear 1, Bear 1, etc. We have written a java code where we have tokenized each word and assigned them a hardcoded value equal to 1.

Reducer Code:

```
public static class Reduce extends
Reducer<Text,IntWritable,Text,IntWritable>
{
public void reduce(Text key, Iterable<IntWritable>
values,Context context)
throws IOException,InterruptedException {
int sum=0;

for(IntWritable x:
values) {
sum+=x.get();
}
context.write(key, new
IntWritable(sum)); }
}
```

We have created a class Reduce which extends class Reducer like that of Mapper.

We define the data types of input and output key/value pair after the class declaration using angle brackets as done for Mapper.

Both the input and the output of the Reducer is a key- value pair.

Input:

The key nothing but those unique words which have been generated after the sorting and shuffling phase: Text

The value is a list of integers corresponding to each key: IntWritable



Example – Bear, [1, 1], etc.

Output:

The key is all the unique words present in the input text file: Text

The value is the number of occurrences of each of the unique words: IntWritable

Example – Bear, 2; Car, 3, etc.

We have aggregated the values present in each of the list corresponding to each key and produced the final answer.

In general, a single reducer is created for each of the unique words, but, you can specify the number of reducer in mapred-site.xml.

Driver Code:

```
Configuration conf= new Configuration();
```

```
Job job = new Job(conf,"My Word Count Program");
```

```
job.setJarByClass(WordCount.class);
```

```
job.setMapperClass(Map.class);
```

```
job.setReducerClass(Reduce.class);
```

```
job.setOutputKeyClass(Text.class);
```

```
job.setOutputValueClass(IntWritable.class);
```

```
job.setInputFormatClass(TextInputFormat.class);
```

```
job.setOutputFormatClass(TextOutputFormat.class);
```

```
Path outputPath = new Path(args[1]);
```

```
//Configuring the input/output path from the filesystem into the
```

```
job FileInputFormat.addInputPath(job, new Path(args[0]));
```

```
FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

In the driver class, we set the configuration of our MapReduce job to run in Hadoop.

We specify the name of the job, the data type of input/output of the mapper and reducer.

We also specify the names of the mapper and reducer classes.

The path of the input and output folder is also specified.

The method setInputFormatClass () is used for specifying that how a Mapper will read the input data or what will be the unit of work. Here, we have chosen



TextInputFormat so that single line is read by the mapper at a time from the input text file.

The main () method is the entry point for the driver. In this method, we instantiate a new Configuration object for the job.

Run the MapReduce code:

The command for running a MapReduce code is:

```
hadoop jar hadoop-mapreduce-example.jar WordCount /  
sample/input /sample/output
```



Experiment – 5

Aim: Write a Map Reduce program that mines weather data. Weather sensors collecting data every hour at many locations across the globe gather a large volume of log data, which is a good candidate for analysis with MapReduce, since it is semi structured and record-oriented.

Weather Report POC-Map Reduce Program to analyse time-temperature statistics and generate report with max/min temperature.

Problem Statement:

1. The system receives temperatures of various cities (Austin, Boston, etc) of USA captured at regular intervals of time on each day in an input file.
2. System will process the input data file and generates a report with Maximum and Minimum temperatures of each day along with time.
3. Generates a separate output report for each city.

Ex: Austin-r-00000

Boston-r-00000

Newjersy-r-00000

Baltimore-r-00000

California-r-00000

Newyork-r-00000

Expected output:- In each output file record should be like this:

25-Jan-2014 Time: 12:34:542 MinTemp: -22.3 Time: 05:12:345 MaxTemp: 35.7

First download input file which contains temperature statistics with time for multiple cities. Schema of record set : CA_25-Jan-2014 00:12:345 15.7 01:19:345 23.1 02:34:542 12.3

CA is city code, here it stands for California followed by date. After that each pair of values represent time and temperature.

Mapper class and map method:-

The very first thing which is required for any map reduce problem is to understand what will be the type of keyIn, ValueIn, KeyOut, ValueOut for the given Mapper class and followed by type of map method parameters.

- `public class WhetherForecastMapper extends Mapper <Object, Text, Text, Text>`
- Object (keyIn) - Offset for each line, line number 1, 2...



- Text (ValueIn) - Whole string for each line (CA_25-Jan-2014 00:12:345)
- Text (KeyOut) - City information with date information as string
- Text (ValueOut) - Temperature and time information which need to be passed to reducer as string.
- public void map(Object keyOffset, Text dayReport, Context con) { } · *KeyOffset* is like line number for each line in input file.
- *dayreport* is input to map method - whole string present in one line of input file.
- *con* is context where we write mapper output and it is used by reducer. **Reducer class and**

Reducer method:-

Similarly, we have to decide what will be the type of keyIn, ValueIn, KeyOut, ValueOut for the given Reducer class and followed by type of reducer method parameters.

- public class WhetherForecastReducer extends Reducer<Text, Text, Text, Text>
- Text(keyIn) - it is same as keyOut of Mapper.
- Text(ValueIn)- it is same as valueOut of Mapper. ·
- Text(KeyOut)- date as string
- text(ValueOut) - reducer writes max and min temperature with time as string
- public void reduce(Text key, Iterable<Text> values, Context context)
- Text key is value of mapper output. i.e:- City & date information
- Iterable<Text> values - values stores multiple temperature values for a given city and date.
- context object is where reducer write it's processed outcome and finally written in file.

Multiple Outputs :- In general, reducer generates output file(i.e: part_r_0000), however in this use case we want to generate multiple output files. In order to deal with such scenario we need to use MultipleOutputs of "org.apache.hadoop.mapreduce.lib.output.MultipleOutputs" which provides a way to write multiple file depending on reducer outcome. See below reducer class for more details. For each reducer task multipleoutput object is created and key/result is written to appropriate file.

Lets create a Map/Reduce project in eclipse and create a class file name it as Calculate Max And Min Temperature With Time. For simplicity, here we have written mapper and reducer class as inner static class. Copy following code lines and paste in newly created class file.

/**



* Question:- To find Max and Min temperature from record set stored in

* text file. Schema of record set :- tab separated (\t) CA_25-Jan-2014

* 00:12:345 15.7 01:19:345 23.1 02:34:542

12.3 03:12:187 16 04:00:093

* -14 05:12:345 35.7 06:19:345 23.1 07:34:542

12.3 08:12:187 16

* 09:00:093 -7 10:12:345 15.7 11:19:345 23.1

12:34:542 -22.3 13:12:187

* 16 14:00:093 -7 15:12:345 15.7 16:19:345

23.1 19:34:542 12.3

* 20:12:187 16 22:00:093 -7

* Expected output:- Creates files for each city and store maximum & minimum

* temperature for each day along with time.

*/

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import
org.apache.hadoop.mapreduce.lib.output.MultipleOutput
s;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.mapreduce.Job;
import
org.apache.hadoop.mapreduce.lib.input.FileInputFormat
;
```



```
import
org.apache.hadoop.mapreduce.lib.output.FileOutputForm
at;
import
org.apache.hadoop.mapreduce.lib.output.TextOutputForm
at;

/**
 * @author devinline
 */

public class CalculateMaxAndMinTemperatureWithTime {
public static String calOutputName =
"California";
public static String nyOutputName = "Newyork";
public static String njOutputName = "Newjersy";
public static String ausOutputName = "Austin";
public static String bosOutputName = "Boston";
public static String balOutputName =
"Baltimore";
public static class WhetherForecastMapper extends
Mapper<Object, Text, Text, Text> {
public void map(Object keyOffset, Text
dayReport, Context con)
throws IOException, InterruptedException {
StringTokenizer strTokens = new
StringTokenizer(
dayReport.toString(), "\t");
int counter = 0;
Float currnetTemp = null;
Float minTemp = Float.MAX_VALUE;
Float maxTemp = Float.MIN_VALUE;
String date = null;
String currentTime = null;
```



```
String minTempANDTime = null;
String maxTempANDTime = null;
while (strTokens.hasMoreElements()) {
    if (counter == 0) {
        date = strTokens.nextToken();
    } else {
        if (counter % 2 == 1) {
            currentTime = strTokens.nextToken();
        } else {
            currnetTemp =
            Float.parseFloat(strTokens.nextToken());
            if (minTemp > currnetTemp) {
                minTemp = currnetTemp;
                minTempANDTime = minTemp + "AND" +
                currentTime;
            }
            if (maxTemp < currnetTemp) {
                maxTemp = currnetTemp;
                maxTempANDTime = maxTemp + "AND" +
                currentTime;
            }
        }
        counter++;
    }
    // Write to context - MinTemp, MaxTemp and
    corresponding time
    Text temp = new Text();
    temp.set(maxTempANDTime);
    Text dateText = new Text();
    dateText.set(date);
    try {
        con.write(dateText, temp);
    }
```

```
} catch (Exception e) {  
    e.printStackTrace();  
}  
temp.set(minTempANDTime);  
dateText.set(date);  
con.write(dateText, temp);  
}  
}
```

public static class WhetherForecastReducer extends

```
Reducer<Text, Text, Text, Text> {  
    MultipleOutputs<Text, Text> mos;  
    public void setup(Context context) {  
        mos = new MultipleOutputs<Text,  
Text>(context);  
    }  
    public void reduce(Text key, Iterable<Text>  
values, Context context)  
throws IOException, InterruptedException {  
        int counter = 0;  
        String reducerInputStr[] = null;  
        String f1Time = "";  
        String f2Time = "";  
        String f1 = "", f2 = "";  
        Text result = new Text();  
        for (Text value : values) {  
            if (counter == 0) {  
                reducerInputStr =  
value.toString().split("AND");  
                f1 = reducerInputStr[0];  
                f1Time = reducerInputStr[1];  
            }  
            else {  
                reducerInputStr =
```



```
value.toString().split("AND");
f2 = reducerInputStr[0];
f2Time = reducerInputStr[1];
}
counter = counter + 1;
}
if (Float.parseFloat(f1) >
Float.parseFloat(f2)) {
result = new Text("Time: " + f2Time + "
MinTemp: " + f2 + "\t"
+ "Time: " + f1Time + " MaxTemp: " + f1);
} else {
result = new Text("Time: " + f1Time + "
MinTemp: " + f1 + "\t"
+ "Time: " + f2Time + " MaxTemp: " + f2);
}
String fileName = "";
if (key.toString().substring(0, 2).equals("CA"))
{
fileName =
CalculateMaxAndMinTemperatureTime.calOutputName;
} else if (key.toString().substring(0,
2).equals("NY")) {
fileName =
CalculateMaxAndMinTemperatureTime.nyOutputName;
} else if (key.toString().substring(0,
2).equals("NJ")) {
fileName =
CalculateMaxAndMinTemperatureTime.njOutputName;
} else if (key.toString().substring(0,
3).equals("AUS")) {
fileName =
CalculateMaxAndMinTemperatureTime.ausOutputName;
```




```
} else if (key.toString().substring(0,
3).equals("BOS")) {
    fileName =
    CalculateMaxAndMinTemperatureTime.bosOutputName;
} else if (key.toString().substring(0,
3).equals("BAL")) {
    fileName =
    CalculateMaxAndMinTemperatureTime.balOutputName;
}
String strArr[] = key.toString().split("_");
key.set(strArr[1]); //Key is date value
mos.write(fileName, key, result);
}
```

@Override

public void cleanup(Context context) **throws**

IOException,

InterruptedException {

mos.close();

}

}

public static void main(String[] args) **throws**

IOException,

ClassNotFoundException,

InterruptedException {

Configuration conf = **new** Configuration();

Job job = Job.getInstance(conf, "Weather

Statistics of USA");

job.setJarByClass(CalculateMaxAndMinTemperatureW
ithTime.class);

job.setMapperClass(WhetherForecastMapper.class);

job.setReducerClass(WhetherForecastReducer.class)

;

job.setMapOutputKeyClass(Text.class);



```
job.setMapOutputValueClass(Text.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(Text.class);
MultipleOutputs.addNamedOutput(job,
    calOutputName,
    TextOutputFormat.class, Text.class,
    Text.class);
MultipleOutputs.addNamedOutput(job,
    nyOutputName,
    TextOutputFormat.class, Text.class,
    Text.class);
MultipleOutputs.addNamedOutput(job,
    njOutputName,
    TextOutputFormat.class, Text.class,
    Text.class);
MultipleOutputs.addNamedOutput(job,
    bosOutputName,
    TextOutputFormat.class, Text.class,
    Text.class);
MultipleOutputs.addNamedOutput(job,
    ausOutputName,
    TextOutputFormat.class, Text.class,
    Text.class);
MultipleOutputs.addNamedOutput(job,
    balOutputName,
    TextOutputFormat.class, Text.class,
    Text.class);
// FileInputFormat.addInputPath(job, new
    Path(args[0]));
// FileOutputFormat.setOutputPath(job, new
    Path(args[1]));
Path pathInput = new Path(
    "hdfs://192.168.213.133:54310/weatherInputData/
```



```
input_temp.txt");
Path pathOutputDir = new Path(
"hdfs://192.168.213.133:54310/user/hduser1/
testfs/output_mapred3");
FileInputFormat.addInputPath(job, pathInput);
FileOutputFormat.setOutputPath(job,
pathOutputDir);
try {
System.exit(job.waitForCompletion(true) ? 0 :
1);
} catch (Exception e) {
// TODO Auto-generated catch block
e.printStackTrace();
}
} }
```

Now execute above sample program. Run -> Run as hadoop. Wait for a moment and check whether output directory is in place on HDFS. Execute following command to verify the same.

```
hduser1@ubuntu:/usr/local/hadoop2.6.1/bin$ ./hadoop
fs -ls /user/hduser1/testfs/output_mapred3
Found 8 items
-rw-r--r-- 3 zytham supergroup 438
2015-12-11 19:21 /user/hduser1/testfs/output_mapred3/
Austin-r-00000
-rw-r--r-- 3 zytham supergroup 219
2015-12-11 19:21 /user/hduser1/testfs/output_mapred3/
Baltimore-r-00000
-rw-r--r-- 3 zytham supergroup 219
2015-12-11 19:21 /user/hduser1/testfs/output_mapred3/
Boston-r-00000
-rw-r--r-- 3 zytham supergroup 511
2015-12-11 19:21 /user/hduser1/testfs/output_mapred3/
```

California-r-00000

-rw-r--r-- 3 zytham supergroup 146

2015-12-11 19:21 /user/hduser1/testfs/output_mapred3/

Newjersy-r-00000

-rw-r--r-- 3 zytham supergroup 219

2015-12-11 19:21 /user/hduser1/testfs/output_mapred3/

Newyork-r-00000

-rw-r--r-- 3 zytham supergroup 0

2015-12-11 19:21 /user/hduser1/testfs/output_mapred3/

_SUCCESS

-rw-r--r-- 3 zytham supergroup 0

2015-12-11 19:21 /user/hduser1/testfs/output_mapred3/

part-r-00000

Open one of the file and verify expected output schema, execute following command for the same.

hduser1@ubuntu:usr/local/hadoop2.6.1/bin\$./hadoop

fs -cat /user/hduser1/testfs/output_mapred3/

Austin-r-00000

25-Jan-2014 Time: 12:34:542 MinTemp: -22.3 Time:

05:12:345 MaxTemp: 35.7

26-Jan-2014 Time: 22:00:093 MinTemp: -27.0 Time:

05:12:345 MaxTemp: 55.7

27-Jan-2014 Time: 02:34:542 MinTemp: -22.3 Time:

05:12:345 MaxTemp: 55.7

29-Jan-2014 Time: 14:00:093 MinTemp: -17.0 Time:

02:34:542 MaxTemp: 62.9

30-Jan-2014 Time: 22:00:093 MinTemp: -27.0 Time:

05:12:345 MaxTemp: 49.2

31-Jan-2014 Time: 14:00:093 MinTemp: -17.0 Time:

03:12:187 MaxTemp: 56.0



Experiment – 6

Aim: Implement Matrix Multiplication with Hadoop Map Reduce.

Implementing Matrix Multiplication with Hadoop Map Reduce

```
import java.io.IOException; import
java.util.*;
import java.util.AbstractMap.SimpleEntry; import
java.util.Map.Entry;
import org.apache.hadoop.fs.Path; import
org.apache.hadoop.conf.*; import
org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat; import
org.apache.hadoop.mapreduce.lib.input.TextInputFormat; import
org.apache.hadoop.mapreduce.lib.output.FileOutputFormat; import
org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
public class TwoStepMatrixMultiplication {
    public static class Map extends Mapper<LongWritable, Text, Text, Text> {
        public void map(LongWritable key, Text value, Context context) throws
IOException, InterruptedException {
            String line = value.toString();
            String[] indicesAndValue = line.split(","); Text outputKey = new
Text();
            Text outputValue = new Text();
            if (indicesAndValue[0].equals("A")) { outputKey.set(indicesAndValue[2]);
                outputValue.set("A," + indicesAndValue[1] + "," +
indicesAndValue[3]);
                context.write(outputKey, outputValue); } else {
                outputKey.set(indicesAndValue[1]); outputValue.set("B," +
indicesAndValue[2] + "," +
indicesAndValue[3]);
                context.write(outputKey, outputValue); }
```



```

    } }

    public static class Reduce extends Reducer<Text, Text, Text, Text> {
        public void reduce(Text key, Iterable<Text> values, Context context) throws
IOException, InterruptedException {
            String[] value;
            ArrayList<Entry<Integer, Float>> listA = new
ArrayList<Entry<Integer, Float>>();
            ArrayList<Entry<Integer, Float>> listB = new
ArrayList<Entry<Integer, Float>>();
            for (Text val : values) {
                value = val.toString().split(","); if
                (value[0].equals("A")) {
                    listA.add(new SimpleEntry<Integer,
Float>(Integer.parseInt(value[1]), Float.parseFloat(value[2])));
                } else {
                    listB.add(new SimpleEntry<Integer,
Float>(Integer.parseInt(value[1]), Float.parseFloat(value[2])));
                } }
            String i; float
            a_ij; String k;
            float b_jk;
            Text outputValue = new Text();
            for (Entry<Integer, Float> a : listA) { i =
                Integer.toString(a.getKey()); a_ij = a.getValue();
                for (Entry<Integer, Float> b : listB) { k =
                    Integer.toString(b.getKey()); b_jk = b.getValue();
                    outputValue.set(i + "," + k + "," +
Float.toString(a_ij*b_jk));
                    context.write(null, outputValue); }
                } }
        }

        public static void main(String[] args) throws Exception { Configuration conf = new
        Configuration();

```



```
        Job job = new Job(conf,
"MatrixMatrixMultiplicationTwoSteps");
        job.setJarByClass(TwoStepMatrixMultiplication.class);
        job.setOutputKeyClass(Text.class); job.setOutputValueClass(Text.class);
        job.setMapperClass(Map.class);
        job.setReducerClass(Reduce.class);
        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputFormatClass(TextOutputFormat.class);
        FileInputFormat.addInputPath(job, new Path("hdfs://
127.0.0.1:9000/matrixin"));
        FileOutputFormat.setOutputPath(job, new Path("hdfs://
127.0.0.1:9000/matrixout"));

        job.waitForCompletion(true);
    }
}
```



Experiment – 7

Aim: Install and Run Pig then write Pig Latin scripts to sort, group, join, project, and filter your data.

Pig Latin scripts to sort, group, join, project, and filter your data.

ORDER BY

Sorts a relation based on one or more fields.

Syntax

```
alias = ORDER alias BY { * [ASC|DESC] | field_alias [ASC|DESC] [, field_alias
[ASC|DESC] ...] } [PARALLEL n];
```

Terms

alias	The name of a relation.
*	The designator for a tuple.
field_ali a s	A field in the relation. The field must be a simple type.
ASC	Sort in ascending order.
DESC	Sort in descending order.
PARALLE	Increase the parallelism of a job by specifying the number of reduce tasks, n.
L n	For more information, see Use the Parallel Features .

Usage

Note: ORDER BY is NOT stable; if multiple records have the same ORDER BY key, the order in which these records are returned is not defined and is not guaranteed to be the same from one run to the next.

In Pig, relations are unordered (see [Relations, Bags, Tuples, Fields](#)):

- If you order relation A to produce relation X (X = ORDER A BY * DESC;) relations A and X still contain the same data.
- If you retrieve relation X (DUMP X;) the data is guaranteed to be in the order you specified (descending).
- However, if you further process relation X (Y = FILTER X BY \$0 > 1;) there is no guarantee that the data will be processed in the order you originally specified (descending).

Pig currently supports ordering on fields with simple types or by tuple designator (*). You cannot order on fields with complex types or by expressions.

```
A = LOAD 'mydata' AS (x: int, y: map[]);
B = ORDER A BY x; -- this is allowed because x is a simple
```




type

B = ORDER A BY y; -- this is not allowed because y is a

complex type

B = ORDER A BY y#'id'; -- this is not allowed because y#'id'

is an expression

Examples

Suppose we have relation A.

A = LOAD 'data' AS (a1:int,a2:int,a3:int);

DUMP A;

(1,2,3)

(4,2,1)

(8,3,4)

(4,3,3)

(7,2,5)

(8,4,3)

In this example relation A is sorted by the third field, f3 in descending order. Note that the order of the three tuples ending in 3 can vary.

X = ORDER A BY a3 DESC;

DUMP X; (7,2,5) (8,3,4) (1,2,3) (4,3,3) (8,4,3) (4,2,1)

Returns each tuple with the rank within a relation.

Syntax

alias = RANK alias [BY { * [ASC|DESC] | field_alias [ASC|DESC] [, field_alias
[ASC|DESC] ...] } [DENSE]];

Terms

alias	The name of a relation.
*	The designator for a tuple.
field_alias	A field in the relation. The field must be a simple type.
ASC	Sort in ascending order.
DESC	Sort in descending order.
DENSE	No gap in the ranking values.



Usage

When specifying no field to sort on, the RANK operator simply prepends a sequential value to each tuple.

Otherwise, the RANK operator uses each field (or set of fields) to sort the relation. The rank of a tuple is one plus the number of different rank values preceding it. If two or more tuples tie on the sorting field values, they will receive the same rank.

NOTE: When using the option **DENSE**, ties do not cause gaps in ranking values.

Examples

Suppose we have relation A.

```
A = load 'data' AS (f1:char array, f2:int, f3:char array);
```

```
DUMP A;
```

```
(David,1,N)
```

```
(Tete,2,N)
```

```
(Ranjit,3,M)
```

```
(Ranjit,3,P)
```

```
(David,4,Q)
```

```
(David,4,Q)
```

```
(Jillian,8,Q)
```

```
(JaePak,7,Q)
```

```
(Michael,8,T)
```

```
(Jillian,8,Q)
```

```
(Jose,10,V)
```

In this example, the RANK operator does not change the order of the relation and simply prepends to each tuple a sequential value.

```
B = rank A;
```

```
dump B;
```

```
(1,David,1,N)
```

```
(2,Tete,2,N)
```

```
(3,Ranjit,3,M)
```

```
(4,Ranjit,3,P)
```



```
(5,David,4,Q)
(6,David,4,Q)
(7,Jillian,8,Q)
(8,JaePak,7,Q)
(9,Michael,8,T)
(10,Jillian,8,Q)
(11,Jose,10,V)
```

In this example, the RANK operator works with f1 and f2 fields, and each one with different sorting order. RANK sorts the relation on these fields and prepends the rank value to each tuple. Otherwise, the RANK operator uses each field (or set of fields) to sort the relation. The rank of a tuple is one plus the number of different rank values preceding it. If two or more tuples tie on the sorting field values, they will receive the same rank.

C = rank A by f1 DESC, f2 ASC;

```
dump C;
(1,Tete,2,N)
(2,Ranjit,3,M)
(2,Ranjit,3,P)
(4,Michael,8,T)
(5,Jose,10,V)
(6,Jillian,8,Q)
(6,Jillian,8,Q)
(8,JaePak,7,Q)
(9,David,1,N)
(10,David,4,Q)
(10,David,4,Q)
```

Same example as previous, but DENSE. In this case there are no gaps in ranking values.

C = rank A by f1 DESC, f2 ASC DENSE;

```
dump C;
```



(1,Tete,2,N)

(2,Ranjit,3,M)

(2,Ranjit,3,P)

(3,Michael,8,T)

(4,Jose,10,V)

(5,Jillian,8,Q)

(5,Jillian,8,Q)

(6,JaePak,7,Q)

(7,David,1,N)

(8,David,4,Q)

(8,David,4,Q)



Experiment – 8

Aim: Install and Run Hive then use Hive to create, alter, and drop databases, tables, views, functions, and indexes.

Hive Databases, Tables, Views, Functions and Indexes

Databases in Hive

The Hive concept of a database is essentially just a *catalog* or *namespace* of tables. However, they are very useful for larger clusters with multiple teams and users, as a way of avoiding table name collisions. It's also common to use databases to organize production tables into logical groups. If you don't specify a database, the default database is used.

The simplest syntax for creating a database is shown in the following example:

```
hive> CREATE DATABASE financials;
```

Hive will throw an error if financials already exists. You can suppress these warnings with this variation:

```
hive> CREATE DATABASE IF NOT EXISTS financials;
```

While normally you might like to be warned if a database of the same name already exists, the IF NOT EXISTS clause is useful for scripts that should create a database on-the-fly, if necessary, before proceeding.

You can also use the keyword SCHEMA instead of DATABASE in all the database-related commands.

At any time, you can see the databases that already exist as follows:

```
hive> SHOW DATABASES;
```

```
default
```

```
financials
```

```
hive> CREATE DATABASE human_resources;
```

```
hive> SHOW DATABASES;
```

```
default
```

```
financials
```

```
human_resources
```

If you have a lot of databases, you can restrict the ones listed using a *regular expression*, a concept we'll explain in LIKE and RLIKE, if it is new to you. The following example lists only those databases that start with the letter h and end with any other characters

(the .* part):



```
hive> SHOW DATABASES LIKE 'h.*';
```

```
human_resources
```

```
hive> ...
```

Hive will create a directory for each database. Tables in that database will be stored in subdirectories of the database directory. The exception is tables in the default database, which doesn't have its own directory.

The database directory is created under a top-level directory specified by the property `hive.metastore.warehouse.dir`, which we discussed in Local Mode Configuration and Distributed and Pseudo distributed Mode Configuration. Assuming you are using the default value for this property, `/user/hive/warehouse`, when the `financials` database is created, Hive will create the directory `/user/hive/warehouse/financials.db`. Note

the `.db` extension.

You can override this default location for the new directory as shown in this example:

```
hive> CREATE DATABASE financials
```

```
> LOCATION '/my/preferred/directory';
```

You can add a descriptive comment to the database, which will be shown by the `DESCRIBE DATABASE <database>` command.

```
hive> CREATE DATABASE financials
```

```
> COMMENT 'Holds all financial tables';
```

```
hive> DESCRIBE DATABASE financials;
```

```
financials    Holds all financial tables
```

```
hdfs://master-server/user/hive/warehouse/financials.db
```

Note that `DESCRIBE DATABASE` also shows the directory location for the database. In this example, the *URI scheme* is `hdfs`. For a MapR installation, it would be `maprfs`. For an Amazon Elastic MapReduce (EMR) cluster, it would also be `hdfs`, but you could set `hive.metastore.warehouse.dir` to use Amazon S3 explicitly (i.e., by specifying `s3n://bucketname/...` as the property value). You could use `s3` as the scheme, but the newer `s3n` is preferred.

In the output of `DESCRIBE DATABASE`, we're showing `master-server` to indicate the *URI authority*, in this case a DNS name and optional port number (i.e., `server:port`) for the “master node” of the file system (i.e., where the *NameNode* service is running for HDFS). If you are running in



pseudo-distributed mode, then the master server will be localhost. For *local* mode, the path will be a local path, `file:///user/hive/warehouse/financials.db`.

If the authority is omitted, Hive uses the master-server name and port defined by the property `fs.default.name` in the Hadoop configuration files, found in the `$HADOOP_HOME/conf` directory. To be clear, `hdfs:///user/hive/warehouse/financials.db` is equivalent to `hdfs://master-server/user/hive/warehouse/financials.db`, where master-server is your master node's DNS name and optional port.

For completeness, when you specify a *relative* path (e.g., *some/relative/path*), Hive will put this under your home directory in the distributed filesystem (e.g., `hdfs:///user/<user-name>`) for HDFS. However, if you are running in *local mode*, your current working directory is used as the parent of *some/relative/path*.

For script portability, it's typical to omit the authority, only specifying it when referring to another distributed filesystem instance (including S3 buckets).

Lastly, you can associate key-value properties with the database, although their only function currently is to provide a way of adding information to the output of `DESCRIBE DATABASE EXTENDED <database>`:

```
hive> CREATE DATABASE financials
      > WITH DBPROPERTIES ('creator' = 'Mark Moneybags', 'date'
= '2012-01-02');
```

```
hive> DESCRIBE DATABASE financials;
financials  hdfs://master-server/user/hive/warehouse/
financials.db
```

```
hive> DESCRIBE DATABASE EXTENDED financials;
financials  hdfs://master-server/user/hive/warehouse/
financials.db
{date=2012-01-02, creator=Mark Moneybags};
```

The `USE` command sets a database as your working database, analogous to changing working directories in a filesystem:

```
hive> USE financials;
```



Now, commands such as `SHOW TABLES;` will list the tables in this database.

Unfortunately, there is no command to show you which database is your current working database!

Fortunately, it's always safe to repeat the `USE ...` command; there is no concept in Hive of nesting of databases.

Recall that we pointed out a useful trick in Variables and Properties for setting a property to print the current database as part of the prompt (Hive v0.8.0 and later):

```
hive> set hive.cli.print.current.db=true;
```

```
hive (financials)> USE default;
```

```
hive (default)> set hive.cli.print.current.db=false;
```

```
hive> ...
```

Finally, you can drop a database:

```
hive> DROP DATABASE IF EXISTS financials;
```

The `IF EXISTS` is optional and suppresses warnings

if `financials` doesn't exist.

By default, Hive won't permit you to drop a database if it contains tables. You can either drop the tables first or append

the `CASCADE` keyword to the command, which will cause the Hive to drop the tables in the database first:

```
hive> DROP DATABASE IF EXISTS financials CASCADE;
```

Using the `RESTRICT` keyword instead of `CASCADE` is equivalent to the default behavior, where existing tables must be dropped before dropping the database. When a database is dropped, its directory is also deleted.

Alter Database

You can set key-value pairs in the `DBPROPERTIES` associated with a database using the `ALTER DATABASE` command. No other metadata about the database can be changed, including its name and directory location:

```
hive> ALTER DATABASE financials SET DBPROPERTIES ('edited-by'  
= 'Joe Db');
```

There is no way to delete or “unset” a `DBPROPERTY`.



Creating Tables

The CREATE TABLE statement follows SQL conventions, but Hive's version offers significant extensions to support a wide range of flexibility where the data files for tables are stored, the formats used, etc. We discussed many of these options in Text File Encoding of Data Values.

In this section, we describe the other options available for the CREATE TABLE statement, adapting the employees table declaration we used previously in Collection Data Types:

```
CREATE TABLE IF NOT EXISTS mydb.employees (
    name          STRING COMMENT 'Employee name',
    salary         FLOAT COMMENT 'Employee salary',
    subordinates ARRAY<STRING> COMMENT 'Names of subordinates',
    deductions    MAP<STRING, FLOAT>
                    COMMENT 'Keys are deductions names, values are
percentages',
    address        STRUCT<street:STRING, city:STRING,
state:STRING, zip:INT>
                    COMMENT 'Home address')
COMMENT 'Description of the table'
TBLPROPERTIES ('creator'='me', 'created_at'='2012-01-02
10:00:00', ...)
LOCATION '/user/hive/warehouse/mydb.db/employees';
```

First, note that you can prefix a database name, mydb in this case, if you're not currently working in the target database.

If you add the option IF NOT EXISTS, Hive will silently ignore the statement if the table already exists. This is useful in scripts that should create a table the first time they run.

However, the clause has a gotcha you should know. If the schema specified differs from the schema in the table that already exists, Hive won't warn you. If your intention is for this table to have the new schema, you'll have to drop the old table, losing your data, and then re-create it. Consider if you should use one or more ALTER TABLE statements to change the existing table schema instead.

See Alter Table for details.



You can add a comment to any column, after the type. Like databases, you can attach a comment to the table itself and you can define one or more table *properties*. In most cases, the primary benefit of TBLPROPERTIES is to add additional documentation in a key-value format. However, when we examine Hive's integration with databases such as DynamoDB (see DynamoDB), we'll see that the TBLPROPERTIES can be used to express essential metadata about the database connection.

Hive automatically adds two table properties: `last_modified_by` holds the username of the last user to modify the table, and `last_modified_time` holds the epoch time in seconds of that modification.

Finally, you can optionally specify a location for the table data (as opposed to *metadata*, which the *metastore* will always hold). In this example, we are showing the default location that Hive would use, `/user/hive/warehouse/mydb.db/employees`, where `/user/hive/warehouse` is the default "warehouse" location (as discussed previously), `mydb.db` is the database directory, and `employees` is the table directory.

By default, Hive always creates the table's directory under the directory for the enclosing database. The exception is the *default* database. It doesn't have a directory under `/user/hive/warehouse`, so a table in the *default* database will have its directory created directly in `/user/hive/warehouse` (unless explicitly overridden).

You can also copy the schema (but not the data) of an existing table:

```
CREATE TABLE IF NOT EXISTS mydb.employees2  
LIKE mydb.employees;
```

This version also accepts the optional `LOCATION` clause, but note that no other properties, including the schema, can be defined; they are determined from the original table.

The `SHOW TABLES` command lists the tables. With no additional arguments, it shows the tables in the current working database. Let's assume we have already created a few other tables, `table1` and `table2`, and we did so in the `mydb` database:

```
hive> USE mydb;
```

```
hive> SHOW TABLES;
```

```
employees
```



```
table1
```

```
table2
```

If we aren't in the same database, we can still list the tables in that database:

```
hive> USE default;
```

```
hive> SHOW TABLES IN mydb;
```

```
employees
```

```
table1
```

```
table2
```

If we have a lot of tables, we can limit the ones listed using a *regular expression*, a concept we'll discuss in detail in LIKE and RLIKE:

```
hive> USE mydb;
```

```
hive> SHOW TABLES 'empl.*';
```

```
employees
```

Not all regular expression features are supported. If you know regular expressions, it's better to test a candidate regular expression to make sure it actually works!

The regular expression in the single quote looks for all tables with names starting with empl and ending with any other characters (the .* part).

We can also use the DESCRIBE EXTENDED mydb.employees command to show details about the table. (We can drop the mydb. prefix if we're currently using the mydb database.) We have reformatted the output for easier reading and we have suppressed many details to focus on the items that interest us now:

```
hive> DESCRIBE EXTENDED mydb.employees;
```

```
name      string  Employee name
```

```
salary    float   Employee salary
```

```
subordinates      array<string>   Names of subordinates
```

```
deductions        map<string,float> Keys are deductions names,
```

```
values are percentages
```



```
address struct<street:string,city:string,state:string,zip:int>
```

```
Home address
```

Detailed **Table** Information

Table(tableName:employees,

dbName:mysql, **owner**:me,

...

location:hdfs://master-server/**user**/hive/warehouse/mysql.db/

employees,

parameters:{creator=me, created_at='2012-01-02 10:00:00',

last_modified_user=me,

last_modified_time=1337544510,

comment:Description of the **table**, ..., ...)

Replacing EXTENDED with FORMATTED provides more readable but also more verbose output.

The first section shows the output of DESCRIBE without EXTENDED or FORMATTED (i.e., the schema including the comments for each column).

If you only want to see the schema for a particular column, append the column to the table name.

Here, EXTENDED adds no additional output:

```
hive> DESCRIBE mysql.employees.salary;
```

```
salary float Employee salary
```

Returning to the extended output, note the line in the description that starts with location:. It shows the full URI path in HDFS to the directory where Hive will keep all the data for this table, as we discussed above.

The tables we have created so far are called *managed* tables or sometimes called *internal* tables, because Hive controls the lifecycle of their data (more or less). As we've seen, Hive stores the data for these tables in a subdirectory under the directory defined by hive.metastore.warehouse.dir (e.g., */user/hive/warehouse*), by default.

When we drop a managed table (see Dropping Tables), Hive deletes the data in the table.



However, managed tables are less convenient for sharing with other tools. For example, suppose we have data that is created and used primarily by *Pig* or other tools, but we want to run some queries against it, but not give Hive *ownership* of the data. We can define an *external* table that points to that data, but doesn't take ownership of it.

External Tables

Suppose we are analyzing data from the stock markets. Periodically, we ingest the data for NASDAQ and the NYSE from a source like Infochimps (<http://infochimps.com/datasets>) and we want to study this data with many tools. (See the data sets named `infochimps_dataset_4777_download_16185` and `infochimps_dataset_4778_download_16677`, respectively, which are actually sourced from Yahoo! Finance.) The schema we'll use next matches the schemas of both these data sources. Let's assume the data files are in the distributed filesystem directory `/data/stocks`.

The following table declaration creates an *external* table that can read all the data files for this comma-delimited data in `/data/stocks`:

```
CREATE EXTERNAL TABLE IF NOT EXISTS stocks (
  exchange      STRING, symbol
                STRING,
  ymd           STRING,
  price_open    FLOAT,
  price_high    FLOAT,
  price_low     FLOAT,
  price_close   FLOAT,
  volume        INT,
  price_adj_close FLOAT)
```

```
ROW FORMAT DELIMITED FIELDS TERMINATED BY ','
```

```
LOCATION '/data/stocks';
```

The **EXTERNAL** keyword tells Hive this table is external and the **LOCATION** ...clause is required to tell Hive where it's located.

Because it's external, Hive does not assume it *owns* the data. Therefore, dropping the table *does not* delete the data, although the *metadata* for the table will be deleted.



There are a few other small differences between managed and external tables, where some HiveQL constructs are not permitted for external tables. We'll discuss those when we come to them.

However, it's important to note that the differences between managed and external tables are smaller than they appear at first. Even for managed tables, you *know* where they are located, so you can use other tools, hadoop dfs commands, etc., to modify and even delete the files in the directories for managed tables. Hive may technically own these directories and files, but it doesn't have full control over them! Recall, in Schema on Read, we said that Hive really has no control over the integrity of the files used for storage and whether or not their contents are consistent with the table schema. Even managed tables don't give us this control.

Still, a general principle of good software design is to express intent. If the data is shared between tools, then creating an external table makes this ownership explicit.

You can tell whether or not a table is managed or external using the output of DESCRIBE EXTENDED tablename. Near the end of the Detailed Table Information output, you will see the following for managed tables:

```
... tableType:MANAGED_TABLE)
```

For external tables, you will see the following:

```
... tableType:EXTERNAL_TABLE)
```

As for managed tables, you can also copy the schema (but not the data) of an existing table: Partitioned, Managed Tables

The general notion of partitioning data is an old one. It can take many forms, but often it's used for distributing load horizontally, moving data physically closer to its most frequent users, and other purposes.

Hive has the notion of partitioned tables. We'll see that they have important performance benefits, and they can help organize data in a logical fashion, such as hierarchically.

We'll discuss partitioned managed tables first. Let's return to our employee stable and imagine that we work for a very large multinational corporation. Our HR people often run



queries with WHERE clauses that restrict the results to a particular country or to a particular *first-level subdivision* (e.g., *state* in the United States or *province* in Canada). (First-level subdivision is an actual term, used here, for example: http://www.commondatahub.com/state_source.jsp.) We'll just use the word *state* for simplicity. We have redundant state information in the address field. It is distinct from the state partition. We could remove the state element from address. There is no ambiguity in queries, since we have to use address.state to project the value inside the address. So, let's partition the data first by country and then by state:

```
CREATE TABLE employees (
  name          STRING,
  salary        FLOAT,
  subordinates  ARRAY<STRING>,
  deductions    MAP<STRING, FLOAT>,
  address       STRUCT<street:STRING, city:STRING,
state:STRING, zip:INT>
)
PARTITIONED BY (country STRING, state STRING);
```

Partitioning tables changes how Hive structures the data storage. If we create this table in the mydb database, there will still be an *employees* directory for the table:

```
hdfs://master_server/user/hive/warehouse/mydb.db/employees
```

However, Hive will now create subdirectories reflecting the

partitioning structure. For example:

```
...
.../employees/country=CA/state=AB
.../employees/country=CA/state=BC
...
.../employees/country=US/state=AL
.../employees/country=US/state=AK
...
```

Yes, those are the actual directory names. The state directories will contain zero or more files for the employees in those states.



Once created, the partition *keys* (country and state, in this case) behave like regular columns. There is one known exception, due to a bug (see Aggregate functions). In fact, users of the table don't need to *care* if these “columns” are partitions or not, except when they want to optimize query performance.

For example, the following query selects all employees in the state of Illinois in the United States:

```
SELECT * FROM employees  
WHERE country = 'US' AND state = 'IL';
```

Note that because the country and state values are encoded in directory names, there is no reason to have this data in the data files themselves. In fact, the data just gets in the way in the files, since you have to account for it in the table schema, and this data wastes space.

Perhaps the most important reason to partition data is for faster queries. In the previous query, which limits the results to employees in Illinois, it is only necessary to scan the contents of *one* directory. Even if we have thousands of country and state directories, all but one can be ignored. For very large data sets, partitioning can dramatically improve query performance, but *only* if the partitioning scheme reflects common *range* filtering (e.g., by locations, timestamp ranges).

When we add predicates to WHERE clauses that filter on partition values, these predicates are called *partition filters*.

Even if you do a query across the entire US, Hive only reads the 65 directories covering the 50 states, 9 territories, and the District of Columbia, and 6 military “states” used by the armed services. You can see the full list here: <http://www.50states.com/abbreviations.htm>.

Of course, if you need to do a query for all employees around the globe, you can still do it. Hive will have to read every directory, but hopefully these broader disk scans will be relatively rare.

However, a query across all partitions could trigger an enormous MapReduce job if the table data and number of partitions are large. A highly suggested safety measure is putting Hive into “strict” mode, which prohibits queries of partitioned tables without a WHERE clause that filters on partitions. You can set the mode to “nonstrict,” as in the following session:



```
hive> set hive.mapred.mode=strict;
```

```
hive> SELECT e.name, e.salary FROM employees e LIMIT 100;
```

FAILED: Error in semantic analysis: No partition predicate

found for

Alias "e" Table "employees"

```
hive> set hive.mapred.mode=nonstrict;
```

```
hive> SELECT e.name, e.salary FROM employees e LIMIT 100;
```

```
John Doe  100000.0
```

...

You can see the partitions that exist with the SHOW

PARTITIONS command:

```
hive> SHOW PARTITIONS employees;
```

...

```
Country=CA/state=AB
```

```
country=CA/state=BC
```

...

```
country=US/state=AL
```

```
country=US/state=AK
```

...

If you have a lot of partitions and you want to see if partitions have been defined for particular partition keys, you can further restrict the command with an optional PARTITION clause that specifies one or more of the partitions with specific values:

```
hive> SHOW PARTITIONS employees PARTITION(country='US');
```

```
country=US/state=AL
```

```
country=US/state=AK
```

...

```
hive> SHOW PARTITIONS employees PARTITION(country='US',
```



```
state='AK');
```

```
country=US/state=AK
```

The DESCRIBE EXTENDED employees command shows the partition keys:

```
hive> DESCRIBE EXTENDED employees;
```

```
name          string,
```

```
salary        float,
```

```
...
```

```
address       struct<...>,
```

```
country       string,
```

```
state         string
```

Detailed **Table** Information...

```
partitionKeys:[FieldSchema(name:country, type:string,
```

```
comment:null),
```

```
FieldSchema (name:state, type:string, comment:null)],
```

```
...
```

The schema part of the output lists the country and state with the other columns, because they are columns as far as queries are concerned. The Detailed Table information includes the Country and State as partition keys. The comments for both of these keys are null; we could have added comments just as for regular columns.

You create partitions in managed tables by loading data into them. The following example creates a US and CA (California) partition while loading data into it from a local directory, *\$HOME/california-employees*. You must specify a value for each partition column. Notice how we reference the HOME environment variable in HiveQL:

```
LOAD DATA LOCAL INPATH '${env:HOME}/california-employees'
```

```
INTO TABLE employees
```

```
PARTITION (country = 'US', state = 'CA');
```

The directory for this partition, *.../employees/country=US/ state=CA*, will be created by Hive and all data files in *\$HOME/california-employees* will be copied into it. See Loading Data into Managed Tables for more information on populating tables.

External Partitioned Tables

You can use partitioning with external tables. In fact, you may find that this is your most common scenario for managing large production data sets. The combination gives you a way to “share”



data with other tools, while still optimizing query performance. You also have more flexibility in the directory structure used, as you define it yourself. We'll see a particularly useful example in a moment.

Let's consider a new example that fits this scenario well: logfile analysis. Most organizations use a standard format for log messages, recording a timestamp, severity (e.g., ERROR, WARNING, INFO), perhaps a server name and process ID, and then an arbitrary text message. Suppose our Extract, Transform, and Load (ETL) process ingests and aggregates logfiles in our environment, converting each log message to a tab-delimited record and also decomposing the timestamp into separate year, month, and day fields, and a combined hms field for the remaining hour, minute, and second parts of the timestamp, for reasons that will become clear in a moment. You could do this parsing of log messages using the string parsing functions built into Hive or Pig, for example. Alternatively, we could use smaller integer types for some of the timestamp-related fields to conserve space. Here, we are ignoring subsequent resolution.

Here's how we might define the corresponding Hive table:

```
CREATE EXTERNAL TABLE IF NOT EXISTS log_messages (  
    hms                INT,  
    severity           STRING,  
    server             STRING  
    process_id         INT,  
    message            STRING)  
PARTITIONED BY (year INT, month INT, day INT)  
ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t';
```

We're assuming that a day's worth of log data is about the correct size for a useful partition and finer grain queries over a day's data will be fast enough.

Recall that when we created the nonpartitioned external stocks table, a `LOCATION ...` clause was required. It isn't used for external partitioned tables. Instead, an `ALTER TABLE` statement is used to add *each* partition separately. It must specify a value for each partition key, the year, month, and day, in this case (see `Alter Table` for more details on this feature). Here is an example, where we add a partition for January 2nd, 2012:

```
ALTER TABLE log_messages ADD PARTITION(year = 2012, month = 1,
```



day = 2)

LOCATION 'hdfs://master_server/data/log_messages/2012/01/02';

The directory convention we use is completely up to us. Here, we

follow a hierarchical directory structure, because it's a logical way to organize our data, but there is no requirement to do so.

You don't have to be an Amazon Elastic MapReduce user to use S3 this way. S3 support is part of the Apache Hadoop distribution. You can *still* query this data, even queries that cross the month-old "boundary," where some data is read from HDFS and some data is read from S3!

By the way, Hive doesn't care if a partition directory doesn't exist for a partition or if it has no files. In both cases, you'll just get no results for a query that filters for the partition. This is convenient when you want to set up partitions before a separate process starts writing data to the. As soon as data is there, queries will return results from that data.

This feature illustrates another benefit: new data can be written to a dedicated directory with a clear distinction from older data in other directories. Also, whether you move old data to an "archive" location or delete it outright, the risk of tampering with newer data is reduced since the data subsets are in separate directories.

As for nonpartitioned external tables, Hive does not own the data and it does not delete the data if the table is dropped.

As for managed partitioned tables, you can see an external table's partitions with **SHOW PARTITIONS**:

```
hive> SHOW PARTITIONS log_messages;
```

...

```
year=2011/month=12/day=31
```

```
    year=2012/month=1/day=1
```

```
    year=2012/month=1/day=2
```

...

Similarly, the **DESCRIBE EXTENDED log_messages** shows the partition keys both as part of the schema and in the list of partitionKeys:



```
hive> DESCRIBE EXTENDED log_messages;
```

```
...
```

```
message          string,
```

```
year            int,
```

```
month          int,
```

```
day            int
```

Detailed **Table** Information...

```
partitionKeys:[FieldSchema(name:year, type:int, comment:null),
```

```
FieldSchema(name:month, type:int, comment:null),
```

```
FieldSchema(name:day, type:int, comment:null)],
```

```
...
```

This output is missing a useful bit of information, the actual location of the partition data. There is a location field, but it only shows Hive's default directory that would be used if the table were a managed table. However, we can get a partition's location as follows:

```
hive> DESCRIBE EXTENDED log_messages PARTITION (year=2012,  
month=1, day=2);
```

```
...
```

```
location:s3n://ourbucket/logs/2011/01/02,
```

```
...
```

We frequently use external partitioned tables because of the many benefits they provide, such as logical data management, performant queries, etc.

ALTER TABLE ... ADD PARTITION is not limited to external tables. You can use it with managed tables, too, when you have (or will have) data for partitions in directories created outside of the LOAD and INSERT options we discussed above. You'll need to remember that not all of the table's data will be under the usual Hive "warehouse" directory, and this data *won't* be deleted when you drop the managed table! Hence, from a "sanity" perspective, it's questionable whether you should dare to use this feature with managed tables.

Customizing Table Storage Formats



In Text File Encoding of Data Values, we discussed that Hive defaults to a text file format, which is indicated by the optional clause `STORED AS TEXTFILE`, and you can overload the default values for the various delimiters when creating the table. Here we repeat the definition of the employees table we used in that discussion:

```
CREATE TABLE employees (  
    name          STRING,  
    salary        FLOAT,  
    subordinates  ARRAY<STRING>,  
    deductions   MAP<STRING, FLOAT>,  
    address       STRUCT<street:STRING, city:STRING,  
state:STRING, zip:INT>  
)
```

ROW FORMAT

FIELDS TERMINATED BY '001'

COLLECTION ITEMS TERMINATED BY '002'

MAP KEYS TERMINATED BY '003'

LINES TERMINATED BY '\n'

STORED AS TEXTFILE;

TEXTFILE implies that all fields are encoded using alphanumeric characters, including those from international character sets, although we observed that Hive uses non-printing characters as “terminators” (delimiters), by default. When TEXTFILE is used, each line is considered a separate record.

You can replace TEXTFILE with one of the other built-in file formats supported by Hive, including SEQUENCEFILE and RCFILE, both of which optimize disk space usage and I/O bandwidth performance using binary encoding and optional compression. Hive draws a distinction between how records are encoded into files and how columns are encoded into records. You customize these behaviors separately.

The record encoding is handled by an *input format* object (e.g., the Java code behind TEXTFILE.) Hive uses a Java *class* (compiled module) named `org.apache.hadoop.mapred.TextInputFormat`. If you are unfamiliar with Java, the dotted name syntax indicates a hierarchical namespace tree



of *packages* that actually corresponds to the directory structure for the Java code. The last name, `TextInputFormat`, is a *class* in the lowest-level package `mapred`.

The record parsing is handled by a *serializer*/

deserializer or *SerDe* for short. For `TEXTFILE` and the encoding we described in Chapter 3 and repeated in the example above, the `SerDe` Hive uses is another Java class called `org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe`.

For completeness, there is also an *output format* that Hive uses for writing the output of queries to files and to the console.

For `TEXTFILE`, the Java class named

`org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat` is used for output.

Third-party input and output formats and SerDes can be specified, a feature which permits users to customize Hive for a wide range of file formats not supported natively.

Here is a complete example that uses a custom SerDe, input format, and output format for files accessible through the *Avro* protocol, which we will discuss in detail in *Avro Hive SerDe*:

```
CREATE TABLE kst
PARTITIONED BY (ds string)
ROW FORMAT SERDE 'com.linkedin.haivvreo.AvroSerDe'
WITH SERDEPROPERTIES ('schema.url'='http://schema_provider/
kst.avsc')
STORED AS
INPUTFORMAT 'com.linkedin.haivvreo.AvroContainerInputFormat'
OUTPUTFORMA
'com.linkedin.haivvreo.AvroContainerOutputFormat';
```

The `ROW FORMAT SERDE ...` specifies the SerDe to use. Hive provides the `WITH SERDEPROPERTIES` feature that allows users to pass configuration information to the SerDe. Hive knows nothing about the meaning of these properties. It's up to the SerDe to decide their meaning. Note that the name and value of each property must be a quoted string.



Finally, the `STORED AS INPUTFORMAT ... OUTPUTFORMAT ...` clause specifies the Java classes to use for the input and output formats, respectively. If you specify one of these formats, you are required to specify both of them.

Note that the `DESCRIBE EXTENDED` table command lists the input and output formats, the SerDe, and any SerDe properties in the `DETAILED TABLE INFORMATION`. For our example, we would see the following:

```
hive> DESCRIBE EXTENDED kst
...
inputFormat:com.linkedin.haivvreo.AvroContainerInputFormat,
outputFormat:com.linkedin.haivvreo.AvroContainerOutputFormat,
...
serdeInfo:SerDeInfo(name:null,
serializationLib:com.linkedin.haivvreo.AvroSerDe,
  parameters:{schema.url=http://schema_provider/kst.avsc})
...
```

Finally, there are a few additional `CREATE TABLE` clauses that describe more details about how the data is supposed to be stored. Let's extend our previous stocks table example from External Tables.

```
CREATE EXTERNAL TABLE IF NOT EXISTS stocks (
  exchange      STRING,
  symbol        STRING,
  ymd           STRING,
  price_open    FLOAT,
  price_high    FLOAT,
  price_low     FLOAT,
  price_close   FLOAT,
  volume        INT,
  price_adj_close FLOAT)
CLUSTERED BY (exchange, symbol)
SORTED BY (ymd ASC)
INTO 96 BUCKETS
ROW FORMAT DELIMITED FIELDS TERMINATED BY ','
```




LOCATION '/data/stocks';

The **CLUSTERED BY ... INTO ... BUCKETS** clause, with an optional **SORTED BY ...** clause is used to optimize certain kinds of queries, which we discuss in detail in Bucketing Table Data Storage.

Dropping Tables: The familiar **DROP TABLE** command from SQL is supported.

DROP TABLE IF EXISTS employees;

The **IF EXISTS** keywords are optional. If not used and the table doesn't exist, Hive returns an error.

For *managed* tables, the table metadata *and* data are deleted.

For *external* tables, the metadata is deleted *but* the data is not.

Alter Table

Most table properties can be altered with **ALTER TABLE** statements, which change *metadata* about the table but not the data itself. These statements can be used to fix mistakes in schema, move partition locations (as we saw in External Partitioned Tables), and do other operations.

Use this statement to rename the table

log_messages to logmsgs:

ALTER TABLE log_messages **RENAME TO** logmsgs;

Adding, Modifying, and Dropping a Table Partition

As we saw previously, **ALTER TABLE** table **ADD PARTITION ...** is used to add a new partition to a table (usually an *external* table). Here we repeat the same command shown previously with the additional options available:

ALTER TABLE log_messages **ADD IF NOT EXISTS**

PARTITION (**year** = 2011, **month** = 1, **day** = 1) **LOCATION** '/logs/

2011/01/01'

PARTITION (**year** = 2011, **month** = 1, **day** = 2) **LOCATION** '/logs/

2011/01/02'

PARTITION (**year** = 2011, **month** = 1, **day** = 3) **LOCATION** '/logs/

2011/01/03'

...;



Multiple partitions can be added in the same query when using Hive v0.8.0 and later. As always, IF NOT EXISTS is optional and has the usual meaning.

Similarly, you can change a partition location, effectively moving it:

```
ALTER TABLE log_messages PARTITION(year = 2011, month = 12,  
day = 2)
```

```
SET LOCATION 's3n://ourbucket/logs/2011/01/02';
```

This command does not move the data from the old location, nor does it delete the old data.

Finally, you can drop a partition:

```
ALTER TABLE log_messages DROP IF EXISTS PARTITION(year = 2011,  
month = 12, day = 2);
```

The IF EXISTS clause is optional, as usual. For managed tables, the data for the partition is *deleted*, along with the metadata, even if the partition was created using ALTER TABLE ... ADD PARTITION. For external tables, the data is not deleted.

There are a few more ALTER statements that affect partitions discussed later in Alter Storage Properties and Miscellaneous Alter Table Statements.

Changing Columns

You can rename a column, change its position, type, or comment:

```
ALTER TABLE log_messages  
CHANGE COLUMN hms hours_minutes_seconds INT  
COMMENT 'The hours, minutes, and seconds part of the  
timestamp'
```

```
AFTER severity;
```

You have to specify the old name, a new name, and the type, even if the name or type is not changing. The keyword COLUMN is optional as is the COMMENT clause. If you aren't moving the column, the AFTER other_column clause is not necessary. In the example shown, we move the column after the severity column. If you want to move the column to the first position, use FIRST instead of AFTER other_column.

As always, this command changes metadata only. If you are moving columns, the data must already match the new schema or you must change it to match by some other means.

Adding Columns



You can add new columns to the end of the existing columns, before any partition columns.

```
ALTER TABLE log_messages ADD COLUMNS (  
  app_name    STRING COMMENT 'Application name',  
  session_id LONG   COMMENT 'The current session id');
```

The COMMENT clauses are optional, as usual. If any of the new columns are in the wrong position, use an ALTER COLUMN table CHANGE COLUMN statement for each one to move it to the correct position.

Deleting or Replacing Columns

The following example removes *all* the existing columns and replaces them with the new columns specified:

```
ALTER TABLE log_messages REPLACE COLUMNS (  
  hours_mins_secs INT           COMMENT 'hour, minute, seconds from  
timestamp',  
  severity                   STRING COMMENT 'The message severity'  
  message                   STRING COMMENT 'The rest of the message');
```

This statement effectively renames the original hms column and removes the server and process_id columns from the original schema definition. As for all ALTER statements, only the table metadata is changed.

The REPLACE statement can only be used with tables that use one of the native *SerDe* modules: DynamicSerDe or MetadataTypedColumn setSerDe. Recall that the SerDe determines how records are parsed into columns (deserialization) and how a record's columns are written to storage (serialization). See Chapter 15 for more details on SerDes.

Alter Table Properties

You can add additional table properties or modify existing properties, but not remove them:

```
ALTER TABLE log_messages SET TBLPROPERTIES (  
  'notes' = 'The process id is no longer captured; this column
```



is always NULL');

Alter Storage Properties

There are several ALTER TABLE statements for modifying format and SerDe properties.

The following statement changes the storage format for a partition to be SEQUENCEFILE, as we discussed in Creating Tables:

```
ALTER TABLE log_messages  
PARTITION(year = 2012, month = 1, day = 1)  
SET FILEFORMAT SEQUENCEFILE;
```

The PARTITION clause is required if the table is partitioned.

You can specify a new SerDe along with SerDe properties or change the properties for the existing SerDe. The following example specifies that a table will use a Java class named com.example.JSONSerDe to process a file of JSON-encoded records:

```
ALTER TABLE table_using_JSON_storage  
SET SERDE 'com.example.JSONSerDe'  
WITH SERDEPROPERTIES (  
  'prop1' = 'value1',  
  'prop2' = 'value2');
```

The SERDEPROPERTIES are passed to the SerDe module (the Java class com.example.JSONSerDe, in this case). Note that both the property names (e.g., prop1) and the values (e.g., value1) must be quoted strings.

The SERDEPROPERTIES feature is a convenient mechanism that SerDe implementations can exploit to permit user customization.

We'll see a real-world example of a JSON SerDe and how it uses SERDEPROPERTIES in JSON SerDe.

The following example demonstrates how to add new SERDEPROPERTIES for the current SerDe:

```
ALTER TABLE table_using_JSON_storage  
SET SERDEPROPERTIES (
```



```
'prop3' = 'value3',  
'prop4' = 'value4');
```

You can alter the storage properties that we discussed in Creating Tables:

```
ALTER TABLE stocks  
CLUSTERED BY (exchange, symbol)  
SORTED BY (symbol)  
INTO 48 BUCKETS;
```

The **SORTED BY** clause is optional, but the **CLUSTER BY** and **INTO ... BUCKETS** are required. (See also Bucketing Table Data Storage for information on the use of data bucketing.)

Miscellaneous Alter Table Statements

In Execution Hooks, we'll discuss a technique for adding execution

“hooks” for various operations. The **ALTER TABLE ... TOUCH** statement is used to trigger these hooks:

```
ALTER TABLE log_messages TOUCH  
PARTITION(year = 2012, month = 1, day = 1);
```

The **PARTITION** clause is required for partitioned tables. A typical scenario for this statement is to trigger execution of the hooks when table storage files have been modified outside of Hive. For example, a script that has just written new files for the 2012/01/01 partition for log_message can make the following call to the Hive CLI:

```
hive -e 'ALTER TABLE log_messages TOUCH PARTITION(year = 2012,  
month = 1, day = 1);'
```

This statement won't create the table or partition if it doesn't already exist. Use the appropriate creation commands in that case.

The **ALTER TABLE ... ARCHIVE PARTITION** statement captures the partition files into a Hadoop archive (HAR) file. This only reduces the number of files in the filesystem, reducing the load on the *NameNode*, but doesn't provide any space savings (e.g., through compression):



```
ALTER TABLE log_messages ARCHIVE  
PARTITION(year = 2012, month = 1, day = 1);
```

To reverse the operation, substitute UNARCHIVE for ARCHIVE. This feature is only available for individual partitions of partitioned tables.

Finally, various protections are available. The following statements prevent the partition from being dropped and queried:

```
ALTER TABLE log_messages  
PARTITION(year = 2012, month = 1, day = 1) ENABLE NO_DROP;
```

```
ALTER TABLE log_messages  
PARTITION(year = 2012, month = 1, day = 1) ENABLE OFFLINE;
```

To reverse either operation, replace ENABLE with DISABLE. These operations also can't be used with nonpartitioned tables.

What is a View?

Views are similar to tables, which are generated based on the requirements.

- We can save any result set data as a view in Hive · Usage is similar to as views used in SQL
- All type of DML operations can be performed on a view Creation of View:

Syntax:

```
Create VIEW < VIEWNAME> AS SELECT
```

Example:

```
Hive>Create VIEW Sample_ViewAS SELECT * FROM employees WHERE  
salary>25000
```

In this example, we are creating view Sample_View where it will display all the row values with salary field greater than 25000.

What is Index?

Indexes are pointers to particular column name of a table.

- The user has to manually define the index
- Wherever we are creating index, it means that we are creating pointer to particular
column name of table



- Any Changes made to the column present in tables are stored using the index value created on the column name.

Syntax:

Create INDEX < INDEX_NAME> ON TABLE <
TABLE_NAME(column names)>

Example:

Create INDEX sample_Index ON TABLE
guruhive_internaltable(id)

Here we are creating index on table guruhive_internaltable for column name id.

HIVE FUNCTIONS

Functions are built for a specific purpose to perform operations like Mathematical, arithmetic, logical and relational on the operands of table column names.

Built-in functions

These are functions that already available in Hive. First, we have to check the application requirement, and then we can use this built in functions in our applications. We can call these functions directly in our application.

The syntax and types are mentioned in the following section.

Types of Built-in Functions in HIVE

- Collection Functions ·
Date Functions
- Mathematical Functions ·
Conditional Functions
- String Functions ·
Misc. Functions

Collection Functions:

These functions are used for collections. Collections mean the grouping of elements and returning single or array of elements depends on return type mentioned in function name.



Return Type	Function Name	Description
INT	size(Map<K.V>)	It will fetch and give the components number in the map type
INT	size(Array<T>)	It will fetch and give the elements number in the array type
Array <K>	Map_keys(Map <K.V>)	It will fetch and gives an array containing the keys of the input map. Here array is in unordered
Array <V>	Map_values(Map <K.V>)	It will fetch and gives an array containing the values of the input map. Here array is in unordered
Array <T>	Sort_array(Array<T>)	sorts the input array in ascending order of array and elements and returns it

Date Functions:

These are used to perform Date Manipulations and Conversion of Date types from one type to another type:

Function Name	Return Type	Description
Unix_Timestamp()	BigInt	We will get current Unix timestamp in seconds
To_date(string timestamp)	string	It will fetch and give the date part of a timestamp string:
year(string date)	INT	It will fetch and give the year part of a date or a timestamp string
quarter(date/ timestamp/string)	INT	It will fetch and give the quarter of the year for a date, timestamp, or string in the range 1 to 4
month(string date)	INT	It will give the month part of a date or a timestamp string
hour(string date)	INT	It will fetch and gives the hour of the timestamp
minute(string date)	INT	It will fetch and gives the minute of the timestamp
Date_sub(string starting date, int days)	string	It will fetch and gives Subtraction of number of days to starting date
Current_date	date	It will fetch and gives the current date at the start of query evaluation
LAST_day(string date)	string	It will fetch and gives the last day of the month which the date belongs to



trunc(string date, string format)	string	It will fetch and gives date truncated to the unit specified by the format. Supported formats in this : MONTH/MON/MM, YEAR/YYYY/ YY.
-----------------------------------	--------	---

Mathematical Functions:

These functions are used for Mathematical Operations. Instead of creating UDFs, we have some inbuilt mathematical functions in Hive.

Function Name	Return Type	Description
round(DOUBLE X)	DOUBLE	It will fetch and returns the rounded BIGINT value of X
round(DOUBLE X, INT d)	DOUBLE	It will fetch and returns X rounded to d decimal places
bround(DOUBLE X)	DOUBLE	It will fetch and returns the rounded BIGINT value of X using HALF_EVEN rounding mode
floor(DOUBLE X)	BIGINT	It will fetch and returns the maximum BIGINT value that is equal to or less than X value
ceil(DOUBLE a), ceiling(DOUBLE a)	BIGINT	It will fetch and returns the minimum BIGINT value that is equal to or greater than X value
rand(), rand(INT seed)	DOUBLE	It will fetch and returns a random number that is distributed uniformly from 0 to 1

Conditional Functions:

These functions used for conditional values checks.

Function Name	Return Type	Description
if(Boolean testCondition, T valueTrue, T valueFalseOrNull)	T	It will fetch and gives value True when Test Condition is of true, gives value False Or Null otherwise.
ISNULL(X)	Boolean	It will fetch and gives true if X is NULL and false otherwise.
ISNOTNULL(X)	Boolean	It will fetch and gives true if X is not NULL and false otherwise.

String Functions:

String manipulations and string operations these functions can be called.